

We address the problem of tractably drawing exact samples from conditional probability distributions. In particular we propose a formalism for constructively incorporating constraints into generative models for probabilistic inference. Two potential implementations are outlined, both of which build upon formalisms of inference as conditional execution of a program, and analyse such programs with respect to their formal semantics.

1.1 INTRODUCTION

Probabilistic inference has established itself as a means of reasoning in domains subject to uncertainty. Although the rules of conditional probability state how to update our beliefs in hypotheses conditioned on evidence, it tells us only declaratively, and leaves us with no guidance as to how hypotheses should be constructed in the first place. In anticipation of the distinction emphasised in this essay, we describe conditional probability and in particular Bayes' Rule as the *logic* of a program, while the *control* or *procedure* remains a topic of continued research.

Inference typically refers to finding the expectation of some function with respect to a probability distribution, computing the posterior distribution, or drawing from it. Numerous approaches exist, but our focus is on Monte Carlo methods whose strengths lie in their applicability to high dimensional distributions, where direct sampling becomes intractable. Inference by means of random perturbations to hypotheses is often inefficient however; we would prefer to not waste resources proposing candidate hypotheses which violate sensible constraints. In other words, we would like to make good guesses.

The objective of this study is to explore conceptual and practical methods of making good guesses by incorporating constraints *constructively* into generative models. We frame construction in contrast to testing; to construct is (ideally) to not test. We build upon recent interpretations of probabilistic inference in computational terms, in particular the `QUERY` operation, which performs universal inference in the probabilistic programming language Church [1]. We suggest a new constructive perspective `CONSTRAIN`, which given a generative model manifested as a stochastic *prior program* P , and a conditioning predicate C aims to construct a new program P^* which samples only values which adhere to our condition.

Universal inference [1] is analagous to notions of universality in computation. Informally it refers to the ability to perform conditional simulation of any computable probabilistic generative process.

The feasibility of this aim is vulnerable to skepticism; how can we construct only constrained proposals without testing to see if they satisfy our constraints? It is unlikely this objective can be achieved in full generality, in some cases we have no alternative than to generate and test. Furthermore we shall see that relaxations of the above maxim may lead to more pragmatic implementation strategies; to construct is to test less perhaps, or to test only partially and incrementally. Precisely what distinguishes one constraint from another in terms of the extent to which it can be implemented constructively, is a topic of great interest and speculated on in our conclusion. Yet, we can motivate the plausibility of constructive inference by appealing to common experience: consider the ease with one can compose a paragraph that rhymes, generate a polynomial expression, or draw an acyclic graph. Each of these examples entails a logical constraint which we satisfy without hinderance. While it is risky to draw conclusions about the computational hardness of a problem based on our subjective difficulty when solving it, each of the above examples has a known efficient algorithm. We conjecture the automatic discovery of algorithms as a means towards conditional sampling, is necessary for any general inference procedure approaching optimal efficiency.

First we will summarise QUERY, a computational theory of probabilistic inference. Then we formalise the key concept of this paper, CONSTRAIN, in terms of QUERY. Two approaches to implementing CONSTRAIN are then described in detail.

1.1.1 Query

QUERY formalises probabilistic inference as a *program* in a corresponding model of computation. Developed first as a primitive function in the probabilistic programming language Church [1], QUERY has been described in terms of probabilistic generalisations of the λ -calculus and Turing machine, both classical models of computation.

In its original formulation, QUERY is a higher order function which accepts as input a stochastic expression to be evaluated and a set of predicate conditions. In lisp (clojure) notation:

*Notation: defn
defines a function.
The following term
is the function name
and terms inside
square brackets []
are argument names.
Let's used to assign
names to values,
conceptually similar
to assignment of
values to variables
in imperative
languages.*

```
(defn query [exp pred]           ; Define function
  (let ((val (eval exp))         ; sample from model, call it val
        (if (pred val)          ; If val satisfies conditions
            val                  ; Then.. return it
            (query exp pred))))  ; Otherwise try again
```

In [3] a formulation in terms of a probabilistic turing machine (PTM) is given. A Turing machine [4] is a mathematical abstraction of a machine, which may read, write and seek access on a finite collection of infinitely long binary tapes. Prior to execution, its input is loaded onto one or more of its tapes, and the output is the content of its

tapes after the machine halts. A probabilistic Turing machine (PTM) is a Turing machine equipped with a tape consisting of a sequence of independent random bits, which is accessible to the Turing machine as a read only randomness source.

QUERY is a PTM which takes two inputs, a prior program P and a conditioning predicate C . Both P and C are themselves encodings of PTMs that take no input. QUERY generates a sample from P . Then, if C is satisfied this sample is outputted, otherwise the process is repeated.

It should be of little surprise that both these formulations are equivalent, as they both perform the function of drawing samples from a prior distribution conditioned on C , using rejection. While rejection sampling provides a simple and intuitive understanding of the meaning of QUERY, it is of course grossly inefficient for the majority of non-trivial problems. Much research in inference is in looking for tractable approximations and alternatives.

1.1.2 *Constrain*

The semantics of CONSTRAIN can be readily understood in terms of QUERY. Expanding on the lisp definition of QUERY given above, CONSTRAIN is a higher order function expecting a prior program P , and taking *constraint* C^* . Uncertain evidence, or equivalently soft constraints are not considered here. Hence a condition C and constraint C^* can be used almost interchangeably, but are differentiated in terms of hardness; constraints are logical and must be true. CONSTRAIN simply returns a function of no arguments which is QUERY with all its arguments evaluated.

```
(defn constrain [exp pred]          ; Define a function
  (fn [] (query expr pred)))       ; Return a 0-ary function
```

While QUERY returns a sample given a model and condition, CONSTRAIN returns a function of no arguments which calls QUERY. In deterministic programs, a function with all its arguments evaluated is simply a value which evaluates to itself, and little is typically gained from taking a functional perspective. In stochastic programs however the output of CONSTRAIN implicitly defines a conditional distribution. This alone differs little from QUERY, as we have only described the semantics. We differentiate with the objective that that C^* is conditioned on constructively; we wish to refrain from applying it as a predicate to fully formed samples, i.e., testing, and instead exploit its semantics to find a more efficient means. The rest of this study is devoted into proposed means of doing this.

CONSTRAINED GENERATIVE MODELS

To recap: we wish to construct a probabilistic program which when evaluated draws samples from a conditional distribution $P^*(x) = P(x|C^*)$. And we aim to do so constructively: to synthesise an efficient procedure of this declarative goal, which possesses improved complexity and efficiency properties.

Possible methods are numerous, and vary principally in dependence on domain specific knowledge. Two proposals will be outlined informally; both draw primarily from subdomains of program semantics: the automated reasoning, formal verification, debugging, and symbolic execution of programs.

2.1 FAULT LOCALISING EVALUATION

Constraining generative models without knowledge specific to the domain seems implausible; how can one generate only convex polygons without any understanding of convexity or geometry? Yet the complexities of representing, incorporating and devising algorithms to exploit domain specific knowledge is best avoided. Black box optimisation methods succeed in obviating domain specific requirements, but at the cost of discarding useful information about the structure of the problem and causes of error in proposals. This approach posits that when considered in accordance to the semantics of program in which they are represented, P , C^* and x are sufficiently meaningful to sidestep any need to inject domain specific knowledge.

The constraint predicate C^* itself contains a wealth of information to aid in both transforming particular samples to force them to adhere to constraints, and transforming the generative model itself. The central idea is of *transparency*: to treat neither constraint C^* nor sample x as monolithic and impenetrable entities. They are viewed instead as structured objects, composed of parts which individually and in concert have a rigorously defined meaning. Failure to satisfy a constraint is not the fault indiscriminately of the whole object then, it can be *blamed* on subcomponents. This idea is easily explained by example, consider the following predicate on a real valued pair of values x_1 and x_2 :

```
(defn C [x1 x2]
  (if (and (pos? x1) (neg? x2))    ; If x1 is +ve and x2 -ve
      true                        ; return true
      false))                    ; else return false
```

Given a sample $(x_1, x_2) = (2.0, 3.0)$ this predicate will return false. An understanding of the program will lead a reasoner to blame this failure on the fact that x_2 is not negative, and furthermore conclude that any change to x_2 to make it negative, such as $x_2 = -0.5$ will satisfy the constraint. Hence in order to fix this particular sample such that it satisfies the constraint, we have found a local constraint that needs to be satisfied, namely x_2 must be negative. We have attributed blame to part of the sample. We can extend this idea to the generative model:

```
(defn gen-pair []
  (list (rand 0 10) (rand -10 10)))
```

Now we can ask see what caused x_2 to be positive, from the generative model we can clearly see that the evaluation of `(rand -10 10)`, which returns a uniform value between -10 and 10, is to blame. It is the cause unsatisfiability at the level of the generative model. In order to correct this, a local constraint can be applied on the second argument of `list` (the cause of the value of x_2). That is, a change to `(rand -10 10)` that will only evaluate to negative values is to be found.

2.1.1.1 An informal algorithm

The concepts of *fault localisation*, *counterfactual reasoning*, *local consistency* and *causal chaining* can be seen in the example, and are central to the proposed method. Informally, the approach can be stated as follows:

1. Evaluate the program until failure (if the constraint is satisfied then we are done), and attribute fault to parts of the program that caused the failure.

A cause of failure could be the evaluation of an expression or the execution of an imperative statement. There may be multiple causes, and interdependencies between them.

2. Consider counterfactuals.

Alternative worlds where a particular cause of failure do not occur are considered. Specifically we seek a local constraint on the cause closest to failure of the entire predicate, that describes the change required result in satisfiability.

3. Chain causes. A causal chain from the cause of failure back to our object of interest is made. This object is either the sample if our goal is to fix one particular sample, or the generative model itself.

4. Transform generative model Taking both our constraints and generative model, we seek a transformation of the model that

causes the constraints to be satisfied. This requires reasoning about the probabilistic semantics of the model.

5. Repeat This algorithm is iterative; we sample a new value from our transformed generative model and continue to refine it with the previous steps.

This approach can be viewed as a special kind of evaluation, where the objective is not to return a value, but instead to work backwards from points of failure to their cause in the generative model. In order to do this we must evaluate both forwards to the point of failure, and backwards to infer the causes.

2.1.2 Program Semantics

This method relies on understanding the meaning of a program, its semantics. To illustrate let us consider a slight variation of the previous predicate:

```
(defn C2 [x1 x2]
  (if (foo (pos? x1) (neg? x2))
      true
      false))
```

We have substituted the *and* for *foo*, a function of unknown semantics. As a result, the program has become meaningless to us. We can no longer reason about its behaviour or infer what the cause of failure will be since *foo* could be any arbitrarily complex or simple function of two arguments. Neither then could we expect any automated reasoner to be able to perform this task.

Fortunately, programming languages are formal objects and their meaning can be specified exactly. Semantics, in a variety of forms (denotational, operational, predicate-transformer, gameplaying) attempts to construct mathematical objects, that describe the meaning of the language. Several practical languages such as Haskell and ML have complete semantics, and are good candidates for use by the proposed method.

2.2 TRANSFORM REJECTIVE GENERATIVE MODEL

Transformational Programming is a prominent method used in automated program development. A formal, declarative specification of a program is *refined* into a complete program by applying a sequence of correctness-preserving transformations. We can appropriate this framework for our purposes; first constrain a naive generative through rejection sampling, then transform into a semantically equivalent, but more efficient program.

From a stochastic program P and constraint C , we construct a new program R_P^C with rejection sampling semantics. R_P^C executes P to sample from its prior and returns the sample C is satisfied, otherwise a further attempt is made. R_P^C is what is returned from application of `CONSTRAIN` to P and C , and can also be viewed as the partial-evaluation of the following lisp function:

Partial evaluation of a program means to take some subset of its arguments, and compile a new program with this subset fixed (under closure) and no longer arguments.

```
(defn R [P C]
  (let [sample (P)]
    (if (true? (C sample))
        sample
        (R P C))))
```

Our next objective is to perform a series of transformations to improve the efficiency R_P^C . By constraining this set of transformations to be semantic preserving - any new program is *equivalent* to the original program, our program will describe the same distribution. It does not follow immediately that a transformed program will be constructive, this depends entirely on the transformations applied. The follow example illustrates one method, which may fall short of constructivism but could improve upon rejection sampling.

Consider a naive generative model which generates polygons by sampling points uniformly over some two dimensional interval. Clearly, the majority of generated polygons will not be convex, most will not even be simple. We wish our polygons to be convex and so specify this as some computable predicate C . R_P^C will construct a fully formed polygon, then reject it if it is not convex. One obvious transformation will result in a new program which applies the convexity test to partially constructed proposals, exploiting the fact that if some part of a polygon is not convex, neither will the whole polygon ever be. Clearly the efficiency gains will depend on the complexity of the test, the size of the polygons being generated and the frequency with which it is applied to partial solutions. Additionally this kind of transformation cannot be applied unconditionally; there instances of constraints which may fail partial solutions but permit whole solutions composed of these failing parts.

2.2.1 Domain General and Specific Transformations

Three well known program transformations are *finite differencing*, *partial evaluation* and *dominated convergence*. Partial evaluation [2] is perhaps the most practical transformation commonly used, and forms the basis of `CONSTRAIN`. It is based on the idea in which a highly parameterised program is concretised by simplification, when some subset of the parameters are fixed. It uses a highly generic strategy which traces a portion of the computation for which expressions can

be evaluated. As a result, partial evaluation must be implemented in full accordance with language semantics

Transformations such as these are deductive in the sense that every one has an associated proof that it is semantic preserving. These proofs are typically domain independent; they apply to any program and exploit the denotations, i.e., the meaning of the program syntax to make valid changes. Clearly however, many semantic preserving transformations will depend on domain specific proofs. A program transformation for the convex polygon example above could depend upon geometric proofs for instance. Automatically generating domain specific transformations then must defer to the problem of automated theorem proving. This is indeed an active area of research in program transformation.

2.3 SAMPLING

First I'll consider a simplified instance of the problem, in anticipation that it will yield insight into the feasibility, or otherwise, of approaches to the more general solution, and may in its own right be applicable to many practical domains.

This simplification can be stated simply as follows: our generative model defines a distribution over linear transformations of a random vector of uniformly distributed real values, where the length of this vector and parameters to each random variable (both lower and upper bound) are fixed and known ahead of time. A constraint involves the conjunction and disjunction of any number of linear equalities and inequalities on these samples. Both the generative model and the constraint are represented as programs in a limited functional language, which importantly is side-effect free and lacks general recursion.

Both a logical and geometric perspective brings some clarity; our constraint program C implicitly defines a logical formula composed of the conjunction and disjunction of a fixed number of literals, each literal representing an inequality of the form $Ax \leq b$ (note we can always convert an equality to two inequalities, as well as convert a greater than relation to a less than through negation). For instance the expression:

```
(if (> x1 10)                ; If val satisfies conditions
    true                    ; then our constraint is satisfied
    (or (> x2 10) (> x1 2))) ; Otherwise if x2>10 or x1>2
```

Defines a logical formula of the form:

$$x_1 > 10 \vee (x_1 \leq 10 \wedge x_2 > 10) \vee (x_1 \leq 10 \wedge x_1 > 2) \quad (1)$$

The extraction of this logical expression from a the program is important, and the first step of the proposed methods. The formula is

one of an infinite number of equivalent alternatives, but is distinct in that it is the disjunction of a number of clauses, where each clause is the conjunction of a number of literals. Each clause can be thought of as a set of local conditions which when all are simultaneously true, will cause the constraint to be satisfied. Geometrically, each literal is a linear inequality defining a half space which splits \mathbb{R}^n into two, and designates one side of the split as feasible and the other infeasible. The entire clause is the intersection of a number of these half spaces and is thus a convex polyhedron. The entire formulae in this form can then be viewed as dividing \mathbb{R}^n into a number of possibly overlapping convex polyhedrons, within which our sample is permitted to lie.

The gist of the proposed method, is to find and isolate these polytopes by statically analysing the predicate, and to focus our sampling within surrounding regions. Restating the objective more formally, we intend to draw samples from $P(X|C)$, where .

Simple Rejection Sampling
Polytope Sampling

2.4 ABSTRACT SAMPLING

Claim: Naive rejection sampling, volume sampling and abstraction sampling all draw samples from the same distribution namely $P(X|C)$.

Definitions $X = (x_1, \dots, x_n)^T$ is a random vector composed of n uniformly distributed random variables. Parameters of these random variables C

2.5 JOINING AND MEETING

In a partially ordered set P , the join and meet of a subset S are respectively the supremum (least upper bound) denoted $\bigvee S$ and infimum (greatest lower bound) denoted $\bigwedge S$, if they exist. If a and b are the elements from P , the join is denoted $a \vee b$ and meet $a \wedge b$.

Join and meet are symmetric duals. In a partially ordered set P gives rise to a dual partially ordered set P^d , which is the set with the inverse order, i.e. $x \geq y$ in P^d iff $y \geq x$ holds in P . Join and meet become symmetric duals with respect to order of inversion.

A lattice is a partially ordered set in which any two elements have a join and a meet.

What is a join here? How does it relate to the lattice join?

A join calculates the topological closure of the convex hull of two polyhedra P_1, P_2 .

There are a few things here: First there is a choice of join. I think a join

2.6 CONCLUSION

Here the probabilistic inference framework `QUERY` was given a constructive perspective with `CONSTRAIN`. We proposed two methods for implementing `CONSTRAIN` based on program transformations and symbolic evaluation. We suspect there will be many difficulties with such an implementation, but preliminary evidence suggests even crude approaches to generating more constrained proposals can have dramatic effects on inference performance.

Discussion has been limited to a single constraint C^* , since multiple constraints can be found conjoining each individual constraint together. However, there may be important practical implications in both of the above methods for the order of the conjunction, e.g. $A \wedge B$ vs $B \wedge A$. There may for instance be a constraint, only visible with a semantic understanding of the programs which represent A and B , such that $A \implies B$. Moreover, A and B may have an internal structure such that there is a more efficient compilation of the two together than just considering each in order.

The central idea of this paper, constructivism, is still open to interpretation. How formally can we differentiate between constructive and non-constructive inference. An appealing possibility is to frame constructivism as the synthesis of an optimal algorithm to sample from the conditional distribution. The cause of variation in the difficulty in which constraints can be adhered to constructively can then be framed in this light; a constraint which is difficult to implement constructively is one where there does not exist (or it is difficult to find) an algorithm which implements it efficiently.

But optimality must be defined with respect to some criterion, and there exist many. An asymptotically optimal algorithm for instance is one which for large inputs performs at worst a constant factor worse than the best possible algorithm. But constant factors can be important, and the worst case is not always indicative of real world performance. Hence, often times even if such an algorithm is known, it is not used regularly.

In spite of these issues the synthesis perspective on inference is appealing, and seeks to appeal to the intuition that whether explicitly or otherwise, almost all approaches to probabilistic inference are implemented on computers as probabilistic programs. The synergy between the fields of program semantics and probability is likely to yield continued insight.

BIBLIOGRAPHY

- [1] N Goodman, V Mansinghka, D Roy, K Bonawitz *Church: a language for generative models* 2012.
- [2] RA Paige, O Danvy *Automatic program development: A tribute to Robert Paige* 2008:
- [3] CE Freer, DM Roy, JB Tenenbaum *Towards common-sense reasoning via conditional simulation: legacies of Turing in Artificial Intelligence* 2012 Arxiv:
- [4] A Turing *On computable numbers, with an application to the Entscheidungsproblem* 1936: Proceedings of the London mathematical society.